


Both collect the knee. The shared queue collects more.



SLO tightness: clock the pinned critical expert needs (MHz)

OLMoE prefill B=16, measured routing; κ from data/kappa_sweep.csv. Work is fixed, so $E = \kappa(f)W$ and only the clock each expert lands on matters. $\kappa(1410)$ is clamped to $\kappa(1380)$ because 1410 throttled, so both savings are understated at the right-hand end.