



prefill B=16

gap: median -3.5%, best layer -5.2%



prefill B=128

gap: median -1.0%, best layer -1.0%



The grey curve IS the existing trade-off: sweeping the single clock buys energy by spending latency, and its minimum (1050 MHz, 1.37x the time) is the cheapest the workload can be run at all. Moving along that curve is NOT the same kind of thing as the coloured curve, which is the same workload with two frequency domains at every latency the grey curve reaches. The vertical GAP is what a second domain actually adds: median 3.5% (B=16) and 1.0% (B=128), and it swings layer to layer because it is a sawtooth in the tile count -- the band is the full spread over 16 layer-samples. Predicted from the fitted model, not measured.