




(a) the split



(b) each partition run alone, swept over clock



(c) kernel cost



For the scheme to save energy the cold partition must prefer a lower clock. It has no preference: both curves are flat within 1% below 1065 MHz and rise together above it. The same holds at 32, 64 and 96 cold experts (appendix). Meanwhile a partitionable kernel costs +5.6% latency before any DVFS.